

2. Linear Regression

March 28, 2022

```
[1]: import pandas as pd
import numpy as np
from sklearn.metrics import mean_squared_error

from sklearn.datasets import fetch_california_housing
X, y = fetch_california_housing(return_X_y=True)

from sklearn.preprocessing import StandardScaler
sc = StandardScaler().fit(X)
X = sc.transform(X)

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

```
[2]: from sklearn.linear_model import LinearRegression
lr = LinearRegression()
lr.fit(X_train, y_train)
y_pred = lr.predict(X_test)
mean_squared_error(y_test, y_pred)
```

```
[2]: 0.5078383014248863
```

```
[3]: # X.shape[1] calculates the number of features in X_train
D = X_train.shape[1]
N = len(X_train) # total number of rows in the dataset
weights = np.random.rand(D) # randomly initialized weight vector
t = y_train # target vector
eta = 0.05 # learning rate
for _ in range(100):
    y = np.dot(X_train, weights)
    error = y - t
    # calculation of OLS loss, error.T will generate transpose
    cost = np.dot(error.T, error)/(2*N)
    #  $w_j = w_j - [(\sum (y-t)(x_j))]/N$ 
    weights = weights - eta * np.dot(X_train.T, error)/N
```

```
y_pred = np.dot(X_test,weights)
mean_squared_error(y_test, y_pred)
```

[3]: 4.906192186434918

[]: